

AI CHEAT SHEET SERIES · COMPLETE PACK

THE COMPLETE *AI MODEL PACK*

All 12 cheat sheets in one PDF — Claude, ChatGPT, Gemini, Llama, Qwen, DeepSeek, Kimi, Hermes, OpenClaw, Grok, Nemotron, and Perplexity. Keep it on your desktop, share it with your team, or print it as a reference guide.

Claude
ANTHROPIC · REASONING

Gemini
GOOGLE · WORKSPACE

Qwen
ALIBABA · LONG CONTEXT

Kimi
MOONSHOT AI · DOCUMENTS

OpenClaw
OPEN SOURCE · AUTOMATION

Nemotron
NVIDIA · ON-PREM

ChatGPT
OPENAI · GENERAL USE

Llama
META · OPEN SOURCE

DeepSeek
DEEPSEEK AI · COST-EFFICIENT

Hermes
NOUS RESEARCH · AGENTIC

Grok
XAI · REAL-TIME

Perplexity
PERPLEXITY AI · SEARCH

Claude

ANTHROPIC

REASONING

“The model that reads, reasons, and writes like your sharpest colleague — careful, thorough, and remarkably good at following complex instructions.”

BEST FOR

- Long-form writing: reports, proposals, contracts, emails
- Deep document analysis and summarization (up to 1M-token context)
- Code review, debugging, and technical documentation
- Nuanced reasoning on ambiguous, multi-step problems
- Safe, compliant deployments in regulated industries

NOT BEST FOR

- ✗ No native image generation (use ChatGPT or Gemini instead)
- ✗ Smaller plugin/app ecosystem than ChatGPT
- ✗ Slightly more conservative on edgy creative content

GETTING STARTED

- Go to [claude.ai](#) and sign up with your email or Google account.
- Start on the Free plan — test it on a real business task before upgrading.
- Try “Projects” to upload business docs Claude can reference in every chat.
- Upgrade to Pro (\$20/mo) once you hit daily usage limits.

PRICING

PLAN	PRICE	NOTES
Free	\$0/mo	Limited daily messages, Sonnet access
Pro	\$20/mo	Higher limits, all models, Projects
Max 5x	\$100/mo	5x Pro usage, priority access
Max 20x	\$200/mo	20x Pro usage, longest sessions
Team	\$30/seat/mo	Shared workspaces, admin controls
Enterprise	Custom	SSO, audit logs, custom contracts

KEY FEATURES

- Models: Opus 5 (deepest reasoning), Sonnet 5 (best balance), Haiku 4.5 (fastest/cheapest), Fable 5 (creative writing)
- Context window: up to 1M tokens on Sonnet 5 and Opus 5; 200K on Haiku 4.5
- Artifacts: live code and document previews inside the chat
- Projects: a persistent knowledge base per workspace
- Claude Code: terminal-based coding agent for full-repo work

PROMPT TEMPLATES

“Rewrite this client email to sound more confident and concise: [paste email]”

“Review this contract clause and flag anything risky or unclear: [paste clause]”

“Summarize this 10-page report into 5 bullet points for my team.”

“Turn these messy meeting notes into a clean action-item list: [paste notes]”

QUICK REFERENCE

Start fresh	New chat, top left
Reference a doc every time	Upload it to a Project
Get code help	Paste the code and describe the bug
Compare options	Ask Claude to build a pros/cons table
Go deeper	Ask “Explain your reasoning step by step”

PRO TIPS

- Use Projects for anything you’ll reference weekly, not just one-off chats.
- Ask Claude to critique its own output for a sharper second draft.
- Long documents: paste the whole thing — no need to summarize first.
- For codebase-wide work, try Claude Code in your terminal.

ChatGPT

OPENAI

GENERAL USE

“The most versatile everyday AI — fast, multimodal, and wired into more tools than anything else on this list.”

BEST FOR

- Everyday tasks: research, writing, brainstorming, planning
- Image generation and editing (Sora, GPT image tools) built in
- Voice mode for hands-free conversations
- Real-time web browsing and Deep Research reports
- The widest plugin, app, and GPT Store ecosystem

NOT BEST FOR

- ✗ Free tier now shows ads on responses (US)
- ✗ The flagship model is gated behind Plus (\$20/mo)
- ✗ Can be verbose without tight, specific prompting

GETTING STARTED

- Go to chatgpt.com and sign up free.
- Try a real business task on the Free plan (GPT-5.3) before upgrading.
- Turn on “Memory” so ChatGPT remembers your business context.
- Upgrade to Plus (\$20/mo) to unlock the GPT-5.6 flagship model, Agent Mode, and Codex.

PRICING

PLAN	PRICE	NOTES
Free	\$0/mo	GPT-5.3, limited use, ads on responses (US)
Go	\$8/mo	Higher limits, no ads, global
Plus	\$20/mo	GPT-5.6 flagship, Agent Mode, Sora, Codex
Pro	\$200/mo	GPT-5.4 Pro, 250 Deep Research runs/mo, 2x context
Business	\$25/seat/mo	Team controls, SOC 2, SSO, 60+ integrations
Enterprise	Custom	Privately hosted, org-wide controls

KEY FEATURES

- Models: GPT-5.6 (Plus flagship), GPT-5.3 (free tier), GPT-5.4 Pro (Pro tier)
- Agent Mode: browses, fills forms, and completes multi-step tasks on its own
- Deep Research: 10 runs/mo on Plus, 250 runs/mo on Pro
- Sora video generation and the Codex coding agent, built in
- Custom GPTs and 60+ Business app integrations

PROMPT TEMPLATES

“Act as a skeptical client and poke holes in this proposal: [paste proposal]”

“Generate 5 social media captions for this offer, each in a different tone: [paste offer]”

“Browse the web and give me the 3 most recent updates on [topic].”

“Turn this rough outline into a polished one-page brief: [paste outline]”

QUICK REFERENCE

Start fresh	New chat, top left
Make an image	Describe it, or upload a reference photo
Get current info	Ask it to search the web first
Automate a task	Switch on Agent Mode and describe the goal
Talk instead of type	Tap the voice icon

PRO TIPS

- Turn on Memory once, then stop repeating your business context every chat.
- Use Deep Research for anything you'd otherwise spend an hour Googling.
- Custom GPTs are worth building for any task you repeat weekly.
- Agent Mode works best with a narrow, well-defined goal, not an open-ended one.

Gemini

GOOGLE

WORKSPACE

“Built into the tools your business already runs on — Gmail, Docs, Sheets — with a context window big enough to read your whole client file.”

BEST FOR

- Native Google Workspace integration (Gmail, Docs, Sheets, Meet)
- Real-time Google Search grounding for current information
- Multimodal input: text, image, video, and audio in one thread
- 1M-token context for research and large-document work
- Deep Research and Canvas for structured, exportable outputs

NOT BEST FOR

- ✗ Free-tier context window is capped at 128K in the app
- ✗ Smaller third-party plugin ecosystem than ChatGPT
- ✗ Best features are gated behind the AI Pro/Ultra tiers

GETTING STARTED

- Go to gemini.google.com and sign in with a Google account.
- The free tier is included with any Google account — test a real task first.
- Connect Gmail, Drive, and Calendar so Gemini can reference your real files.
- Upgrade to Google AI Pro (\$19.99/mo) for Gemini 3.1 Pro's full 1M-token context.

PRICING

PLAN	PRICE	NOTES
Free	\$0/mo	Gemini in-app, 128K context, basic features
AI Plus	\$7.99/mo	128K context, 200 monthly credits, 400GB storage
AI Pro	\$19.99/mo	Gemini 3.1 Pro, full 1M context, 1,000 credits
AI Ultra	\$99.99+/mo	Highest limits, earliest access to new models

KEY FEATURES

- Models: Gemini 3.1 Pro (deep reasoning, 1M context), Gemini 3.5 Flash (fast, agentic/coding)
- Workspace extensions: drafts in Gmail, analyzes Sheets, summarizes Docs live
- 1,000 monthly AI credits on Pro for Flow, Whisk, and NotebookLM
- 400GB+ shared cloud storage across Gmail, Drive, and Photos (Plus tier)
- Real-time grounding in Google Search results

PROMPT TEMPLATES

“Draft a reply to this email in my Gmail thread, keep it under 100 words: [paste thread]”

“Read this Google Sheet of leads and flag anything that looks like a duplicate.”

“Search for the latest news on [topic] and summarize the top 3 developments.”

“Turn this Doc into a one-page client-facing summary with headers.”

QUICK REFERENCE

Start fresh	New chat, top left
Work inside Gmail/Docs/Sheets	Open the Gemini side panel
Get current info	It's grounded in Search by default
Analyze a big file	Upload it directly — 1M-token context on Pro
Build a report	Use Canvas for a structured, editable output

PRO TIPS

- Turn on the Gmail/Docs/Sheets side panel — most of the value is in-context, not the standalone app.
- Upgrade to AI Pro before you paste large documents — the free tier truncates at 128K.
- Use Deep Research for anything that needs cited, current sources.
- NotebookLM (included in credits) is worth a look for turning documents into study guides or audio briefs.

Llama

META

OPEN SOURCE

“Free to run, free to modify, free to own — the model of choice when the data can't leave the building.”

BEST FOR

- Privacy-first deployments — run entirely on your own infrastructure
- Custom fine-tuning on proprietary business data
- Cost-efficient on-prem or private-cloud AI at scale
- Massive-context document processing (Scout: 10M-token context)
- Multimodal apps needing an open-weight foundation model

NOT BEST FOR

- ✗ No official Meta-hosted business chat app — you (or a vendor) host it
- ✗ Requires technical setup or a hosting partner
- ✗ Support is community-driven, not a vendor helpdesk

GETTING STARTED

- Try it free first at meta.ai — Meta's own consumer app, no setup required.
- For business use, pick a hosting route: self-host, or a managed API (Together AI, Groq, AWS Bedrock, Azure).
- Download weights directly from llama.com or Hugging Face if self-hosting.
- Start with the smaller model in a family before scaling to the largest.

PRICING

PLAN	PRICE	NOTES
Weights	Free	Download and run anywhere under the stated license
Self-hosted	Your infra cost	Full control and privacy; requires GPU capacity
Managed API	Pay-per-token	Via Together AI, Groq, Bedrock, etc. — no infra to manage
Meta AI app	Free	Consumer chat interface, no business features

KEY FEATURES

- Llama 4 Maverick: 400B-parameter mixture-of-experts model, natively multimodal (text + image)
- Llama 4 Scout: 10M-token context window — the largest of any open-weight model
- Mixture-of-experts architecture: high capability at lower serving cost than dense models
- Open license: free to self-host, fine-tune, and deploy commercially
- 40+ Fortune 500 companies already run production Llama deployments

PROMPT TEMPLATES

“Fine-tune candidate: draft 10 example Q&A pairs from our support docs for training data: [paste docs]”

“Summarize this internal policy doc in plain language for new hires: [paste doc]”

“Given this customer transcript, extract the 3 main complaints as a list: [paste transcript]”

“Write a system prompt for a self-hosted assistant that only answers from our knowledge base.”

QUICK REFERENCE

Just try it, no setup	meta.ai (consumer app)
Need it private/on-prem	Self-host the weights on your own servers
Need it fast, no infra	Use a managed API provider
Need it to know your data	Fine-tune on your documents
Need huge documents	Use Llama 4 Scout's 10M-token context

PRO TIPS

- If you don't have GPU infrastructure, start with a managed API — it's still the open-weight model, just hosted.
- Match model size to the task: don't run the 400B model for something a smaller checkpoint handles fine.
- Fine-tuning pays off fastest on narrow, repetitive tasks (support replies, intake forms).
- Confirm your hosting partner's data-retention policy before sending client data.

Qwen

ALIBABA

LONG CONTEXT

“Bilingual, budget-friendly, and built for the highest-volume workloads — a trillion-parameter model priced like a mid-tier one.”

BEST FOR

- Bilingual (Chinese/English) and multilingual business tasks
- 1M+ token context windows across the whole model family
- Budget-sensitive, high-volume deployments
- Multimodal input: text, image, and video in one prompt (Qwen3.8-Max)
- Coding and agentic workflows at competitive pricing

NOT BEST FOR

- ✗ The newest flagship (Qwen3.8-Max) is still rolling out to all regions
- ✗ Documentation and support skew toward developers, not businesses
- ✗ Confirm data handling and compliance for regulated industries

GETTING STARTED

- Go to qwen.ai (or the Alibaba Cloud app) and sign up free.
- Test the free chat tier on a real business task.
- For API access, create an Alibaba Cloud Model Studio account.
- Choose your tier: Plus for cost, Max for top capability.

PRICING

PLAN	PRICE	NOTES
Free chat	\$0/mo	qwen.ai web/app, rate-limited
Qwen3.7 Plus (API)	\$0.40 / \$1.60 per 1M tok	Best cost-to-capability ratio
Qwen3.8-Max (API)	\$2.00 / \$6.00 per 1M tok	Top-tier reasoning and multimodal
Alibaba Cloud Enterprise	Custom	SLAs, dedicated capacity, compliance support

KEY FEATURES

- Qwen3.8-Max: 2.4T total / 95B active parameters, 1M context, multimodal (Aug 2026)
- Qwen3.7 Plus: same 1M context at roughly 6x lower cost than Max
- Strong bilingual reasoning and translation quality
- Open-weight variants available for self-hosting alongside the closed flagship
- Fast release cadence — new checkpoints roughly every 4-6 weeks

PROMPT TEMPLATES

“Translate this client email to Mandarin, keep the tone professional but warm: [paste email]”

“Read this 200-page vendor contract and list every clause that mentions liability: [paste/upload]”

“Write product descriptions in English and Chinese for this item: [paste details]”

“Given this dataset, identify the 3 biggest cost-saving opportunities: [paste data]”

QUICK REFERENCE

Start fresh	New chat on qwen.ai
Work bilingually	Just switch languages mid-conversation
Process a huge file	Upload directly — 1M-token context
Keep costs down	Use Qwen3.7 Plus instead of Max
Build a product	Use the Model Studio API

PRO TIPS

- Default to the Plus tier unless a task specifically needs Max-level reasoning — the cost gap is roughly 5x.
- Qwen's bilingual quality makes it a strong pick for any business with Chinese-speaking clients or vendors.
- Batch high-volume, repetitive API calls together to make the most of the low per-token pricing.
- Watch for new checkpoints — this family updates faster than most competitors.

DeepSeek

DEEPSEEK AI

COST-EFFICIENT

“Frontier-level reasoning and coding at a fraction of the cost — and you can download the weights and run it yourself.”

BEST FOR

- Frontier-level reasoning and coding at industry-low pricing
- Long-horizon agent workflows (1M-token context, 384K max output)
- Open-weight self-hosting for full data control
- High-volume API usage where per-token cost compounds fast
- Math, knowledge, and software-engineering benchmark performance

NOT BEST FOR

- ✗ No official “R2” release yet — ignore rumors, use V4 for now
- ✗ Interface is more developer-facing than a polished consumer app
- ✗ Data residency is in China — confirm suitability for regulated data

GETTING STARTED

- Go to chat.deepseek.com or the mobile app (iOS/Android) — free.
- Toggle on “Web Search” per conversation if you need current information.
- Toggle on “DeepThink” (R1-style step-by-step reasoning) for harder problems.
- For production use, use the API with cached-prompt pricing to cut costs further.

PRICING

PLAN	PRICE	NOTES
Free chat	\$0/mo	chat.deepseek.com, rate-limited
API (cache miss)	\$0.435 / \$0.87 per 1M tok	V4 Pro, standard requests
API (cache hit)	\$0.0036 / \$0.87 per 1M tok	Repeated/system prompts
Self-hosted	Free (weights)	Requires your own GPU infrastructure

KEY FEATURES

- DeepSeek-V4 Pro: 1.6T-parameter mixture-of-experts model, deepest reasoning tier
- DeepSeek-V4 Flash: 284B-parameter MoE, faster and cheaper for everyday use
- Context window: 1M tokens by default, up to 384K max output
- Cache-hit input pricing near-zero for repeated or system prompts
- Open weights available — self-host for full privacy

PROMPT TEMPLATES

- “Debug this Python function – explain the bug before fixing it: [paste code]”
- “Walk through this math/finance problem step by step, showing your reasoning: [paste problem]”
- “Search the web and summarize the latest developments on [topic] in 3 bullets.”
- “Refactor this SQL query for performance and explain each change: [paste query]”

QUICK REFERENCE

Start fresh	New chat, top left
Get current info	Turn on Web Search for that message
See the reasoning	Turn on DeepThink (R1-style)
Cut API costs	Reuse the same system prompt to hit the cache
Keep it private	Self-host the open weights

PRO TIPS

- DeepThink mode is worth the extra latency for anything with multiple correct-looking wrong answers (math, logic, code).
- If you're calling the API repeatedly with the same system prompt, structure it to maximize cache hits — the savings are significant.
- Treat any “R2” announcement as a rumor until DeepSeek ships an official model card.
- For sensitive data, self-hosting the open weights is the safest route.



MOONSHOT AI

DOCUMENTS

“Built for the documents that break other models — 200k+ token contracts, research papers, and codebases, handled in one pass.”

BEST FOR

- Processing 200K+ token documents: contracts, academic papers, research
- Long-horizon coding across Python, Rust, and Go
- Converting prompts and screenshots into production-ready UI code
- Multi-agent orchestration for complex, multi-step tasks
- Cost-efficient mixture-of-experts architecture

NOT BEST FOR

- ✗ Smaller context ceiling (262K) than Qwen, DeepSeek, or Gemini’s 1M
- ✗ Less brand recognition and support infrastructure outside Asia
- ✗ Best coding results come from the specialized K2.7 Code variant

GETTING STARTED

- Go to [kimi.com](#) and create a free account.
- Upload the long document you want analyzed — contracts, papers, reports.
- For coding tasks, ask for the K2.7 Code model specifically if available.
- Use the Moonshot AI platform API for production integrations.

PRICING

PLAN	PRICE	NOTES
Free chat	\$0/mo	kimi.com, rate-limited
K2.6 (API, blended)	~\$1.15–\$2.15 per 1M tok	Across tracked providers
K2.7 Code (API)	\$0.68 / \$3.40 per 1M tok	Coding-specialized pricing
Enterprise	Custom	Moonshot AI platform, dedicated support

KEY FEATURES

- Kimi K2.6: 1T total / 32B active parameter MoE, multimodal, flagship model
- Kimi K2.7 Code: specialized end-to-end coding model, same 262K context
- Context window: 262,144 tokens in and out
- Strong at converting visual mockups into working front-end code
- Multi-agent orchestration for chaining complex tasks

PROMPT TEMPLATES

“Read this 150-page contract and list every clause that could hurt us: [upload/paste]”

“Build a working landing page from this screenshot: [attach image]”

“Review this pull request across all changed files and flag risks: [paste diff]”

“Summarize this academic paper’s methodology and findings in plain English: [upload]”

QUICK REFERENCE

Start fresh	New chat on kimi.com
Analyze a huge document	Upload it directly — 262K-token context
Turn a mockup into code	Attach the screenshot and describe the stack
Get better code results	Ask for the K2.7 Code model
Chain multiple steps	Describe the full workflow, let it orchestrate

PRO TIPS

- For contract or research review, paste the whole document — don’t pre-summarize, you’ll lose detail Kimi would have caught.
- Use K2.7 Code specifically for coding tasks; the base K2.6 model is tuned more broadly.
- Screenshot-to-code works best with a clear, single-screen mockup rather than a full multi-page flow.
- If your document exceeds 262K tokens, split it and summarize each part before combining.

Hermes

NOUS RESEARCH

AGENTIC

“Not just a model — a self-improving open-source agent that builds and refines its own skills as it works.”

BEST FOR

- Autonomous agents that run long, multi-step jobs unattended
- Function calling and reliable tool-use across 20+ platforms
- Persistent memory and a learning loop that improves with use
- Agent-to-agent communication for multi-agent systems
- Privacy-first, self-hosted automation (fully open-source, MIT license)

NOT BEST FOR

- Not a simple chat app — expect some technical setup
- Best results require pairing it with an underlying LLM you provide
- Younger ecosystem than the major consumer assistants

GETTING STARTED

- Go to hermes-agent.org or the NousResearch/hermes-agent repo on GitHub.
- Install Hermes Desktop (Mac, Windows, or Linux) for a native GUI — no CLI required.
- Or spin up a hosted instance on Nous Portal Cloud, billed hourly from credits.
- Connect your model provider (or a local model) and your first tools/skills.

PRICING

PLAN	PRICE	NOTES
Self-hosted	Free (MIT license)	Full control; your own compute + model cost
Hermes Desktop	Free	Native GUI app for Mac, Windows, Linux
Nous Portal Cloud	Pay-as-you-go	Hosted instances billed hourly from credits
Underlying model cost	Varies	Pair with Claude, GPT, Llama, or a local model

KEY FEATURES

- v0.20.0 “The Herald Release” (Aug 2026): real-time voice with streaming text-to-speech
- Agent-to-agent communication (A2A v1.0) for multi-agent workflows
- Signed outbound webhooks and grounded research with verifiable citations
- Learning loop: creates, tests, and refines its own skills from experience
- Fully open-source under the MIT license, runs on 20+ platforms

PROMPT TEMPLATES

- “Build me a skill that checks our inbox every morning and drafts replies to routine questions.”
- “Monitor this webpage/RSS feed and alert me by webhook when [condition] happens.”
- “Research [topic] and produce a cited brief, then save it to my notes.”
- “Coordinate with [another agent] to handle intake, then hand off qualified leads to me.”

QUICK REFERENCE

Get a GUI, no terminal	Install Hermes Desktop
Skip the setup entirely	Use a hosted instance on Nous Portal Cloud
Give it a new capability	Ask it to build/refine a skill — it learns from use
Connect two agents	Use A2A v1.0 messaging
Get talked to, not just texted	Use real-time voice mode

PRO TIPS

- Start with Hermes Desktop before touching the CLI — it covers most day-to-day use.
- Let the learning loop run for a week before judging it; skills improve with repeated use.
- Pair it with a cheaper model for routine tasks and a stronger one for anything high-stakes.
- Review signed webhooks and permissions before connecting it to anything business-critical.

OpenClaw

OPEN SOURCE

AUTOMATION

“The personal AI agent that lives on your machine — reads your files, works your apps, and runs while you sleep.”

BEST FOR

- Targeted automation across your own files, apps, and messaging tools
- Connecting AI directly to WhatsApp, Discord, and local file systems
- Model-agnostic setups — bring your own API key or run a local model
- Repetitive, high-volume operational tasks run unattended
- Privacy-first automation that runs on your own Mac, Windows, or Linux machine

NOT BEST FOR

- ✗ Requires comfort installing and configuring local software
- ✗ No single official vendor — support is community-driven (Discord/GitHub)
- ✗ You're responsible for the underlying model's API costs

GETTING STARTED

- Install OpenClaw on Mac, Windows, or Linux from the official GitHub repo.
- Connect an API key from any supported provider — or run a fully local model.
- Enable AgentSkills for the tools you need (100+ preconfigured: shell, files, web).
- Connect WhatsApp, Discord, or your file system to start automating real tasks.

PRICING

PLAN	PRICE	NOTES
Software	Free (open-source)	Self-hosted, install from GitHub
Model/API cost	Varies by provider	Bring your own key (or run local, free)
Local-only setup	Free	No cloud API cost using local models via Ollama

KEY FEATURES

- Fully open-source and model-agnostic — swap providers at runtime without rebuilding
- 100+ preconfigured AgentSkills: shell commands, file management, web automation
- Supports Anthropic, OpenAI, Google, DeepSeek, Moonshot AI, and more, plus local models via Ollama
- Persistent memory across conversations and sessions
- Runs entirely on your own machine — nothing leaves it unless you route it out

PROMPT TEMPLATES

“Watch this folder and file every new invoice into the right client subfolder.”

“Every morning, check WhatsApp for unread client messages and draft replies for my review.”

“Scrape [site] daily and alert me in Discord when a new listing matches [criteria].”

“Rename and organize every file in this download folder by type and date.”

QUICK REFERENCE

First install	Clone the GitHub repo, follow the setup guide
Pick a brain	Add any provider's API key, or point it at a local model
Give it a job	Turn on the relevant AgentSkill(s)
Connect messaging	Link WhatsApp or Discord in settings
Run it unattended	Schedule the task, let it work in the background

PRO TIPS

- Start with a single, low-risk automation before connecting it to anything client-facing.
- Use a cheap/local model for routine automations and reserve a frontier model for judgment calls.
- Review exactly which AgentSkills are enabled — only turn on what a given automation needs.
- Because it's community-run, check the GitHub repo directly for the current release before installing.

Grok

X AI

REAL-TIME

“Wired straight into X for real-time signal — built for speed, edge, and questions that need to know what's happening right now.”

BEST FOR

- Real-time information pulled from X (Twitter) and the live web
- Fast, less-filtered answers than most consumer assistants
- Long-running agentic tasks (Grok 4.6's focus area)
- Multi-agent reasoning on hard problems (Grok 4 Heavy)
- Image understanding alongside text

NOT BEST FOR

- ✗ Smaller context window (500K) than Gemini, DeepSeek, or Qwen's 1M
- ✗ Top tiers (\$100-\$300/mo) are expensive relative to competitors
- ✗ Less enterprise compliance tooling than Claude or ChatGPT

GETTING STARTED

- Go to [grok.com](#) or open Grok inside the X app.
- The free tier is included with any X account — test a real task first.
- Upgrade to SuperGrok (\$30/mo) for full Grok 4 access.
- Step up to Heavy (\$300/mo) only if you need multi-agent parallel reasoning.

PRICING

PLAN	PRICE	NOTES
Free	\$0/mo	Limited access, X account required
X Premium	\$8/mo	Grok bundled with an X subscription
SuperGrok	\$30/mo	Full Grok 4 access, 128K context
Plus	\$100/mo	Higher limits, priority access
Heavy	\$300/mo	Grok 4 Heavy multi-agent reasoning, 256K context

KEY FEATURES

- Grok 4.6 (Aug 2026): 500K-token context, long-running agents, image + text input
- Grok 4 Heavy: multi-agent reasoning with parallel test-time compute, top benchmark tier
- Native, real-time access to X posts and trends — no separate search toggle
- Function calling and structured outputs for building on the API
- Voice and video modes for hands-free or visual tasks

PROMPT TEMPLATES

“What is everyone on X saying about [topic] in the last 24 hours?”

“Give me the unfiltered take on this business decision: [paste decision]”

“Analyze this image and tell me what's actually going on: [attach image]”

“Run this as a long agent task: monitor [topic] and summarize it daily.”

QUICK REFERENCE

Start fresh	New chat on grok.com or in the X app
Get real-time info	Just ask — X data is native, no toggle needed
Analyze an image	Attach it directly
Run a long task	Use an agent task, not a single chat turn
Need deepest reasoning	Use the Heavy tier for multi-agent mode

PRO TIPS

- Grok's edge is live X data — lean on it for anything time-sensitive a static model would miss.
- Don't pay for Heavy unless you actually have multi-step problems that benefit from parallel agents.
- Cross-check anything opinion-flavored — Grok is intentionally less filtered than most competitors.
- For coding or business tasks without a real-time angle, SuperGrok is usually enough.

Nemotron

NVIDIA

ON - PREM

“Nvidia's open-weight engine for agentic AI you run yourself — distilled to fit on a single GPU without giving up frontier-level reasoning.”

BEST FOR

- High-volume agentic workloads that need to run on your own hardware
- Single-GPU deployment — runs on a laptop or desktop, not just a data center
- Post-training on your own data via Nvidia's NeMo framework
- 1M-token context for large-scale document or codebase work
- Enterprise deployments already built on Nvidia infrastructure

NOT BEST FOR

- No consumer chat app — this is a developer/enterprise model, not a website you sign into
- Requires ML engineering resources to deploy and fine-tune
- Best results need Nvidia's own tooling (NeMo, NIM) wrapped around it

GETTING STARTED

- Download Nemotron 3.5 Lightning from Hugging Face or Nvidia's website.
- Run it locally — it's sized to fit on a single GPU.
- Use Nvidia's NeMo framework to post-train it on your own data.
- For production, deploy via Nvidia NIM microservices for a hosted-style API.

PRICING

PLAN	PRICE	NOTES
Weights	Free (open-source)	Download from Hugging Face or Nvidia
Self-hosted	Your infra cost	Runs on a single GPU
Nvidia NIM	Pay-per-use	Hosted microservice deployment
Enterprise support	Custom	Nvidia AI Enterprise licensing

KEY FEATURES

- Nemotron 3.5 Lightning (Aug 2026): 30B parameters, 3B active per token, distilled from Nemotron 3 Ultra
- Context window: up to 1M tokens
- Open-source weights, free to self-host and fine-tune
- Nemotron 4 (in development): flagship targeting 1T+ parameters, expected late 2026
- Already in production at CrowdStrike (security), Harvey (legal), and CodeRabbit (code review)

PROMPT TEMPLATES

“Fine-tune candidate: generate 20 example support tickets and resolutions from our docs: [paste docs]”

“Build an agent that triages incoming security alerts by severity: [describe alerts]”

“Summarize this 500-page technical manual into a quick-reference guide.”

“Review this codebase for common vulnerability patterns: [paste code]”

QUICK REFERENCE

Just try the weights	Hugging Face or Nvidia's model page
Need it on one GPU	Nemotron 3.5 Lightning is built for that
Need it to know your data	Post-train with NeMo
Need a hosted API	Deploy via Nvidia NIM
Need frontier scale	Wait for Nemotron 4 (targeting late 2026)

PRO TIPS

- This is the pick when “open-weight” isn't enough and you specifically need Nvidia's deployment stack (NeMo, NIM) around the model.
- Lightning's single-GPU footprint makes it worth testing before reaching for a bigger, pricier model.
- Treat Nemotron 4 as roadmap, not available — Lightning and the Nemotron 3 family are what's actually shipped today.
- Pair it with NeMo Switchyard if you're routing across multiple models and want cost-aware selection automatically.

Perplexity

PERPLEXITY AI

SEARCH

“Built to answer, not just chat — a research engine that cites its sources and now browses and acts for you through Comet.”

BEST FOR

- Research questions that need current, cited sources
- Quick fact-checking with linked evidence instead of a bare answer
- Comet browser agent: form fills, bookings, email drafting while you do something else
- Comparing answers across top models (GPT-5, Claude Opus) in one place
- Multimedia answers with generated images and video

NOT BEST FOR

- Not built for long, open-ended creative writing or brainstorming
- Max tier (\$200/mo) only pays off if you actually use the autonomous agent features
- Less model customization/fine-tuning than developer-first platforms

GETTING STARTED

- Go to perplexity.ai and ask a real question — no account needed for basic search.
- Sign up free to save threads and get more daily searches.
- Install the Comet browser (free) to get agent features built into browsing.
- Upgrade to Pro (\$20/mo) for extended search limits and access to top-tier models.

PRICING

PLAN	PRICE	NOTES
Free	\$0/mo	Basic search, limited daily queries
Education Pro	\$10/mo	Student-verified discount on Pro
Pro	\$20/mo	Extended search, top-tier models, Comet, file uploads
Max	\$200/mo	Highest limits, autonomous agent tasks run in the background
Enterprise Pro	\$40/seat/mo	Team admin controls

KEY FEATURES

- Every answer comes with cited, clickable sources by default
- Comet: an agentic browser that fills forms, books, drafts email, and shops for you
- Pro gives access to GPT-5, Claude Opus, and other frontier models in one interface
- Comet Plus (\$5/mo, included in Pro/Max): access to premium publisher content
- Image and video generation built directly into search results

PROMPT TEMPLATES

“Research [topic] and give me a cited summary with the 3 most credible sources.”

“Compare these two vendors and cite where each claim comes from: [names]”

“Book me a table for 4 at 7pm Friday somewhere highly rated nearby.”

“Track this product's price and email me when it drops below \$[X].”

QUICK REFERENCE

Start fresh	New thread, top left
Get cited answers	Just ask — sources are automatic
Automate a task	Use the Comet browser agent
Compare models	Switch the model dropdown mid-thread on Pro
Go deep on a topic	Use Pro Search / Deep Research mode

PRO TIPS

- Perplexity's real value is the citations — always check the linked sources before trusting a sensitive claim.
- Comet is worth installing even on the free plan; the agent features are what separate it from a normal search engine.
- Switch models mid-conversation on Pro if one model's answer style doesn't fit the task.
- Don't upgrade to Max unless you'll actually delegate real tasks to the background agent — Pro covers most research needs.