

# DeepSeek

DEEPSEEK AI

COST-EFFICIENT

“Frontier-level reasoning and coding at a fraction of the cost — and you can download the weights and run it yourself.”

## BEST FOR

- Frontier-level reasoning and coding at industry-low pricing
- Long-horizon agent workflows (1M-token context, 384K max output)
- Open-weight self-hosting for full data control
- High-volume API usage where per-token cost compounds fast
- Math, knowledge, and software-engineering benchmark performance

## NOT BEST FOR

- ✗ No official “R2” release yet — ignore rumors, use V4 for now
- ✗ Interface is more developer-facing than a polished consumer app
- ✗ Data residency is in China — confirm suitability for regulated data

## GETTING STARTED

- Go to chat.deepseek.com or the mobile app (iOS/Android) — free.
- Toggle on “Web Search” per conversation if you need current information.
- Toggle on “DeepThink” (R1-style step-by-step reasoning) for harder problems.
- For production use, use the API with cached-prompt pricing to cut costs further.

## PRICING

| PLAN             | PRICE                        | NOTES                                |
|------------------|------------------------------|--------------------------------------|
| Free chat        | \$0/mo                       | chat.deepseek.com, rate-limited      |
| API (cache miss) | \$0.435 / \$0.87 per 1M tok  | V4 Pro, standard requests            |
| API (cache hit)  | \$0.0036 / \$0.87 per 1M tok | Repeated/system prompts              |
| Self-hosted      | Free (weights)               | Requires your own GPU infrastructure |

## KEY FEATURES

- DeepSeek-V4 Pro: 1.6T-parameter mixture-of-experts model, deepest reasoning tier
- DeepSeek-V4 Flash: 284B-parameter MoE, faster and cheaper for everyday use
- Context window: 1M tokens by default, up to 384K max output
- Cache-hit input pricing near-zero for repeated or system prompts
- Open weights available — self-host for full privacy

## PROMPT TEMPLATES

- “Debug this Python function – explain the bug before fixing it: [paste code]”
- “Walk through this math/finance problem step by step, showing your reasoning: [paste problem]”
- “Search the web and summarize the latest developments on [topic] in 3 bullets.”
- “Refactor this SQL query for performance and explain each change: [paste query]”

## QUICK REFERENCE

|                   |                                               |
|-------------------|-----------------------------------------------|
| Start fresh       | New chat, top left                            |
| Get current info  | Turn on Web Search for that message           |
| See the reasoning | Turn on DeepThink (R1-style)                  |
| Cut API costs     | Reuse the same system prompt to hit the cache |
| Keep it private   | Self-host the open weights                    |

## PRO TIPS

- DeepThink mode is worth the extra latency for anything with multiple correct-looking wrong answers (math, logic, code).
- If you're calling the API repeatedly with the same system prompt, structure it to maximize cache hits — the savings are significant.
- Treat any “R2” announcement as a rumor until DeepSeek ships an official model card.
- For sensitive data, self-hosting the open weights is the safest route.