

Llama

META

OPEN SOURCE

“Free to run, free to modify, free to own — the model of choice when the data can't leave the building.”

BEST FOR

- Privacy-first deployments — run entirely on your own infrastructure
- Custom fine-tuning on proprietary business data
- Cost-efficient on-prem or private-cloud AI at scale
- Massive-context document processing (Scout: 10M-token context)
- Multimodal apps needing an open-weight foundation model

NOT BEST FOR

- ✗ No official Meta-hosted business chat app — you (or a vendor) host it
- ✗ Requires technical setup or a hosting partner
- ✗ Support is community-driven, not a vendor helpdesk

GETTING STARTED

- Try it free first at meta.ai — Meta's own consumer app, no setup required.
- For business use, pick a hosting route: self-host, or a managed API (Together AI, Groq, AWS Bedrock, Azure).
- Download weights directly from llama.com or Hugging Face if self-hosting.
- Start with the smaller model in a family before scaling to the largest.

PRICING

PLAN	PRICE	NOTES
Weights	Free	Download and run anywhere under the stated license
Self-hosted	Your infra cost	Full control and privacy; requires GPU capacity
Managed API	Pay-per-token	Via Together AI, Groq, Bedrock, etc. — no infra to manage
Meta AI app	Free	Consumer chat interface, no business features

KEY FEATURES

- Llama 4 Maverick: 400B-parameter mixture-of-experts model, natively multimodal (text + image)
- Llama 4 Scout: 10M-token context window — the largest of any open-weight model
- Mixture-of-experts architecture: high capability at lower serving cost than dense models
- Open license: free to self-host, fine-tune, and deploy commercially
- 40+ Fortune 500 companies already run production Llama deployments

PROMPT TEMPLATES

“Fine-tune candidate: draft 10 example Q&A pairs from our support docs for training data: [paste docs]”

“Summarize this internal policy doc in plain language for new hires: [paste doc]”

“Given this customer transcript, extract the 3 main complaints as a list: [paste transcript]”

“Write a system prompt for a self-hosted assistant that only answers from our knowledge base.”

QUICK REFERENCE

Just try it, no setup	meta.ai (consumer app)
Need it private/on-prem	Self-host the weights on your own servers
Need it fast, no infra	Use a managed API provider
Need it to know your data	Fine-tune on your documents
Need huge documents	Use Llama 4 Scout's 10M-token context

PRO TIPS

- If you don't have GPU infrastructure, start with a managed API — it's still the open-weight model, just hosted.
- Match model size to the task: don't run the 400B model for something a smaller checkpoint handles fine.
- Fine-tuning pays off fastest on narrow, repetitive tasks (support replies, intake forms).
- Confirm your hosting partner's data-retention policy before sending client data.