

# Llama

META

OPEN SOURCE

“Free to run, free to modify, free to own — the model of choice when the data can't leave the building.”

## BEST FOR

- Privacy-first deployments — run entirely on your own infrastructure
- Custom fine-tuning on proprietary business data
- Cost-efficient on-prem or private-cloud AI at scale
- Massive-context document processing (Scout: 10M-token context)
- Multimodal apps needing an open-weight foundation model

## NOT BEST FOR

- ✗ No official Meta-hosted business chat app — you (or a vendor) host it
- ✗ Requires technical setup or a hosting partner
- ✗ Support is community-driven, not a vendor helpdesk

## GETTING STARTED

- Try it free first at meta.ai — Meta's own consumer app, no setup required.
- For business use, pick a hosting route: self-host, or a managed API (Together AI, Groq, AWS Bedrock, Azure).
- Download weights directly from llama.com or Hugging Face if self-hosting.
- Start with the smaller model in a family before scaling to the largest.

## PRICING

| PLAN        | PRICE           | NOTES                                                     |
|-------------|-----------------|-----------------------------------------------------------|
| Weights     | Free            | Download and run anywhere under the stated license        |
| Self-hosted | Your infra cost | Full control and privacy; requires GPU capacity           |
| Managed API | Pay-per-token   | Via Together AI, Groq, Bedrock, etc. — no infra to manage |
| Meta AI app | Free            | Consumer chat interface, no business features             |

## KEY FEATURES

- Llama 4 Maverick: 400B-parameter mixture-of-experts model, natively multimodal (text + image)
- Llama 4 Scout: 10M-token context window — the largest of any open-weight model
- Mixture-of-experts architecture: high capability at lower serving cost than dense models
- Open license: free to self-host, fine-tune, and deploy commercially
- 40+ Fortune 500 companies already run production Llama deployments

## PROMPT TEMPLATES

- “Fine-tune candidate: draft 10 example Q&A pairs from our support docs for training data: [paste docs]”
- “Summarize this internal policy doc in plain language for new hires: [paste doc]”
- “Given this customer transcript, extract the 3 main complaints as a list: [paste transcript]”
- “Write a system prompt for a self-hosted assistant that only answers from our knowledge base.”

## QUICK REFERENCE

|                           |                                           |
|---------------------------|-------------------------------------------|
| Just try it, no setup     | meta.ai (consumer app)                    |
| Need it private/on-prem   | Self-host the weights on your own servers |
| Need it fast, no infra    | Use a managed API provider                |
| Need it to know your data | Fine-tune on your documents               |
| Need huge documents       | Use Llama 4 Scout's 10M-token context     |

## PRO TIPS

- If you don't have GPU infrastructure, start with a managed API — it's still the open-weight model, just hosted.
- Match model size to the task: don't run the 400B model for something a smaller checkpoint handles fine.
- Fine-tuning pays off fastest on narrow, repetitive tasks (support replies, intake forms).
- Confirm your hosting partner's data-retention policy before sending client data.