

Nemotron

NVIDIA

ON - PREM

“Nvidia's open-weight engine for agentic AI you run yourself — distilled to fit on a single GPU without giving up frontier-level reasoning.”

BEST FOR

- High-volume agentic workloads that need to run on your own hardware
- Single-GPU deployment — runs on a laptop or desktop, not just a data center
- Post-training on your own data via Nvidia's NeMo framework
- 1M-token context for large-scale document or codebase work
- Enterprise deployments already built on Nvidia infrastructure

NOT BEST FOR

- ✗ No consumer chat app — this is a developer/enterprise model, not a website you sign into
- ✗ Requires ML engineering resources to deploy and fine-tune
- ✗ Best results need Nvidia's own tooling (NeMo, NIM) wrapped around it

GETTING STARTED

- Download Nemotron 3.5 Lightning from Hugging Face or Nvidia's website.
- Run it locally — it's sized to fit on a single GPU.
- Use Nvidia's NeMo framework to post-train it on your own data.
- For production, deploy via Nvidia NIM microservices for a hosted-style API.

PRICING

PLAN	PRICE	NOTES
Weights	Free (open-source)	Download from Hugging Face or Nvidia
Self-hosted	Your infra cost	Runs on a single GPU
Nvidia NIM	Pay-per-use	Hosted microservice deployment
Enterprise support	Custom	Nvidia AI Enterprise licensing

KEY FEATURES

- Nemotron 3.5 Lightning (Aug 2026): 30B parameters, 3B active per token, distilled from Nemotron 3 Ultra
- Context window: up to 1M tokens
- Open-source weights, free to self-host and fine-tune
- Nemotron 4 (in development): flagship targeting 1T+ parameters, expected late 2026
- Already in production at CrowdStrike (security), Harvey (legal), and CodeRabbit (code review)

PROMPT TEMPLATES

“Fine-tune candidate: generate 20 example support tickets and resolutions from our docs: [paste docs]”

“Build an agent that triages incoming security alerts by severity: [describe alerts]”

“Summarize this 500-page technical manual into a quick-reference guide.”

“Review this codebase for common vulnerability patterns: [paste code]”

QUICK REFERENCE

Just try the weights	Hugging Face or Nvidia's model page
Need it on one GPU	Nemotron 3.5 Lightning is built for that
Need it to know your data	Post-train with NeMo
Need a hosted API	Deploy via Nvidia NIM
Need frontier scale	Wait for Nemotron 4 (targeting late 2026)

PRO TIPS

- This is the pick when “open-weight” isn't enough and you specifically need Nvidia's deployment stack (NeMo, NIM) around the model.
- Lightning's single-GPU footprint makes it worth testing before reaching for a bigger, pricier model.
- Treat Nemotron 4 as roadmap, not available — Lightning and the Nemotron 3 family are what's actually shipped today.
- Pair it with NeMo Switchyard if you're routing across multiple models and want cost-aware selection automatically.